

Pulse – Module Extraction AI Agent

By- CHERRY SHARMA (RA2211042010051)

Technical Architecture, Approach, Assumptions & Edge Case Handling

1. Overview

The **Pulse – Module Extraction AI Agent** is an AI-powered application designed to automatically extract **structured product intelligence** from documentation-based help websites.

Modern SaaS products expose functionality through large and often fragmented documentation portals. For Product Managers, understanding the **functional breakdown of a product (modules and submodules)** is time-consuming and error-prone when done manually.

This project addresses that problem by building a **Generative AI-driven system** that:

- Accepts one or more documentation URLs
- Crawls and processes relevant documentation pages
- Infers modules (major product areas)
- Infers submodules (specific functionalities under each module)
- Generates accurate, PM-style descriptions
- Outputs results in a structured JSON format suitable for downstream consumption

The solution is built as a **local full-stack GenAI application**, focusing on reasoning quality, structure, and maintainability rather than UI complexity.

2. Technical Architecture Description

2.1 High-Level Architecture

The system follows a **layered architecture** with clear separation of responsibilities:

User Interface (React / Streamlit-like UI)

|

| REST API Call

v

Backend Orchestration Layer (FastAPI)

```
|  
|-- URL Validation  
|-- Documentation Crawler  
|-- Content Cleaner & Normalizer  
|-- LLM-based Module Inference Engine
```

|

v

Structured JSON Output

This architecture ensures:

- **Modularity**
- **Ease of debugging**
- **Independent evolution of crawling, AI logic, and UI layers**

2.2 Component-Level Breakdown

2.2.1 Input Layer (Frontend / UI)

Responsibilities

- **Accept one or more documentation URLs**
- **Trigger module extraction**
- **Display structured results**

Design Decisions

- **Simple UI with a single responsibility: enable interaction with the AI agent**
- **No authentication or persistence to keep scope aligned with assignment**
- **Output rendered in a human-readable hierarchical format**

2.2.2 Backend API Layer

Technology

- **Python + FastAPI**

Responsibilities

- **Validate incoming URL inputs**
- **Coordinate crawling, processing, and AI inference**
- **Enforce structured output schema**

Key API Contract

POST /extract

Input: { urls: [string] }

Output: { modules: [...] }

FastAPI was chosen for:

- **Strong request/response validation**
- **Clean API semantics**
- **Production-grade performance characteristics**

2.2.3 Documentation Crawler

Responsibilities

- **Recursively crawl documentation pages starting from the input URL(s)**
- **Follow only internal documentation links**
- **Handle redirects, broken links, and unsupported pages gracefully**

Content Filtering

The crawler removes:

- **Headers**
- **Footers**
- **Navigation menus**
- **Scripts and styles**

This ensures that only **meaningful documentation content** reaches the AI layer.

Controls Implemented

- **Maximum page crawl limit**
- **Timeout handling**
- **Duplicate URL avoidance**

These controls prevent runaway crawling and noise accumulation.

2.2.4 Content Processing & Normalization

Responsibilities

- **Convert raw HTML into clean, normalized text**
- **Preserve logical structure (headings, lists) where possible**
- **Truncate content intelligently to respect LLM token limits**

This layer ensures the AI receives **signal-rich input**, improving accuracy and consistency.

2.2.5 Module & Submodule Inference (GenAI Layer)

This is the **core intelligence** of the system.

Responsibilities

- Infer top-level product modules from documentation sections
- Group related features as submodules
- Generate concise, accurate descriptions for both levels

Key Characteristics

- LLM is used strictly for reasoning and inference
- All outputs are derived only from extracted documentation content
- A strict JSON schema is enforced to avoid ambiguity

Why GenAI?

Traditional rule-based or keyword-based approaches fail because:

- Documentation varies widely in structure
- Similar features are described using different terminology
- Product hierarchies require semantic understanding

LLMs excel at:

- Semantic grouping
- Contextual inference
- Hierarchical reasoning

3. Detailed Approach

3.1 End-to-End Flow

1. **Input Acquisition**
 - User provides one or more documentation URLs
2. **Validation**
 - URLs are validated for format and accessibility
3. **Recursive Crawling**
 - Relevant internal documentation pages are collected
4. **Content Cleaning**
 - Non-informational HTML elements are removed
5. **Normalization**
 - Text is flattened and prepared for AI consumption
6. **LLM-Based Inference**
 - Modules, submodules, and descriptions are inferred
7. **Structured Output**
 - Results returned as structured JSON

3.2 Mapping to Assignment Requirements

Assignment Requirement	How It Is Addressed
Accept one or more URLs	Supported via API/UI
Recursive crawling	Implemented with link traversal
Content filtering	Headers, footers, nav removed
Hierarchy inference	LLM-based semantic reasoning
Structured JSON output	Strict schema enforced
PM-style descriptions	Prompt-guided generation

4. Assumptions

The implementation is based on the following assumptions:

1. Documentation URLs are publicly accessible
2. Documentation content accurately reflects product features
3. HTML structure is reasonably well-formed
4. LLM has sufficient context to infer hierarchy
5. The goal is product understanding, not legal or contractual accuracy

These assumptions are realistic for internal product intelligence tools used by PMs.

5. Edge Case Handling

5.1 Broken Links & Redirects

- Skipped gracefully
- Crawling continues for remaining pages
- Partial but valid results returned

5.2 Deeply Nested Documentation

- Crawl limits prevent infinite traversal
- AI inference relies on semantic grouping rather than strict depth

5.3 Sparse or Poorly Structured Content

- LLM infers structure from semantics instead of headings alone
- Still produces usable modules where possible

5.4 Unsupported or Non-Documentation URLs

- Input validation prevents crashes
- Empty or minimal output returned safely

5.5 LLM Output Formatting Errors

- Strict JSON parsing enforced
- Failures handled gracefully without server crash

6. Best Practices Followed

Engineering Best Practices

- Modular architecture
- Single responsibility per file
- Clean function naming
- Schema validation with Pydantic

GenAI Best Practices

- Low-temperature inference
- Explicit structured prompts
- Token-aware input truncation
- No hallucinated content injection

Reliability & Safety

- Environment variables for secrets
- Graceful failure handling
- Defensive parsing and validation

7. Limitations

- No support for authenticated documentation
- No caching mechanism in current version
- No confidence scoring per module
- No persistence or database layer

These are **intentional MVP trade-offs**, not architectural constraints.

8. Testing & Validation

The system was tested on multiple real-world documentation sources, including:

- <https://support.neo.space/hc/en-us>
- [Instagram](#)
- [X. It's what's happening](#)
- <https://wordpress.org/documentation/>
- <https://help.zluri.com/>
- <https://www.chargebee.com/docs/2.0/>

Results demonstrated:

- **Consistent module grouping**
- **Meaningful submodule extraction**
- **High-quality PM-style descriptions**

9. Conclusion

The Module Extraction AI Agent demonstrates:

- **Strong alignment with Product Management workflows**
- **Practical, high-signal use of Generative AI**
- **Clean, extensible full-stack architecture**
- **Ability to deliver a meaningful MVP within tight timelines**

This project closely aligns with Pulse's mission of **revolutionizing product management using AI**, and can serve as a foundation for more advanced product intelligence capabilities.